



Understanding GANs



Brickki vs. The Forgers: The Cat & Mouse Game of AI



Your First Day – Trouble Already

First day as **Head of Quality Control**.

Company: **Brickki** – maker of high-quality building blocks.

Issue: counterfeit bricks mixed into our product bags.



The Problem - Counterfeit Invasion



Fake bricks **damage customer trust.**

Must be detected before reaching customers.

Your mission: intercept fakes on the production line.

Becoming the Expert Detector

Initial detection is **tricky**.

Customer **feedback** helps identify patterns.

Gradually become skilled at spotting subtle defects.



The Forgers Strike Back

Forgers **improve** fake brick quality.

Differences are now minimal.

Detection becomes much harder.

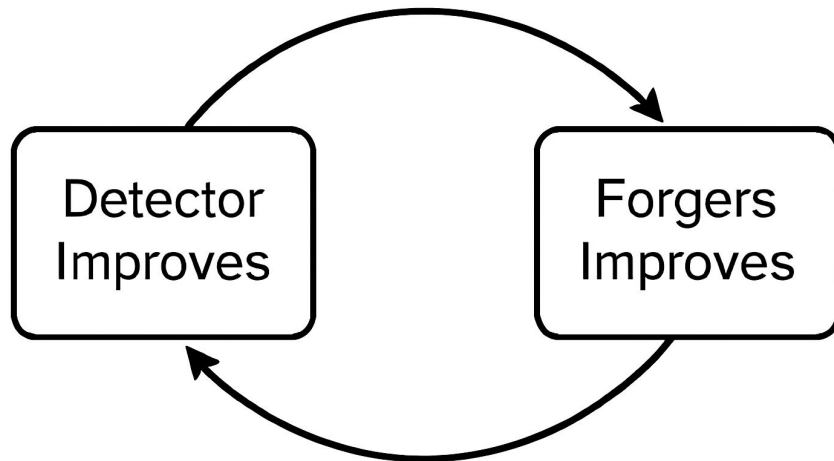


Retaliation & Retraining

You **study** new fakes and improve your detection skills.

Forgers adapt their methods again.

Cycle repeats: both sides improving.




The Endless Cat & Mouse Game

Both real and fake bricks get better over time.

Differences become extremely subtle.

Rivalry drives **continuous improvement**.



From Bricks to AI

Forgers → Generator (creates fakes).

You → Discriminator (detects fakes).

Both improve by competing.

What Makes GANs So Impressive?

Can create photo-realistic images, videos, audio.

Used for:

- Image generation
- Super-resolution
- Deepfakes
- Art & design




2014



2015



2016

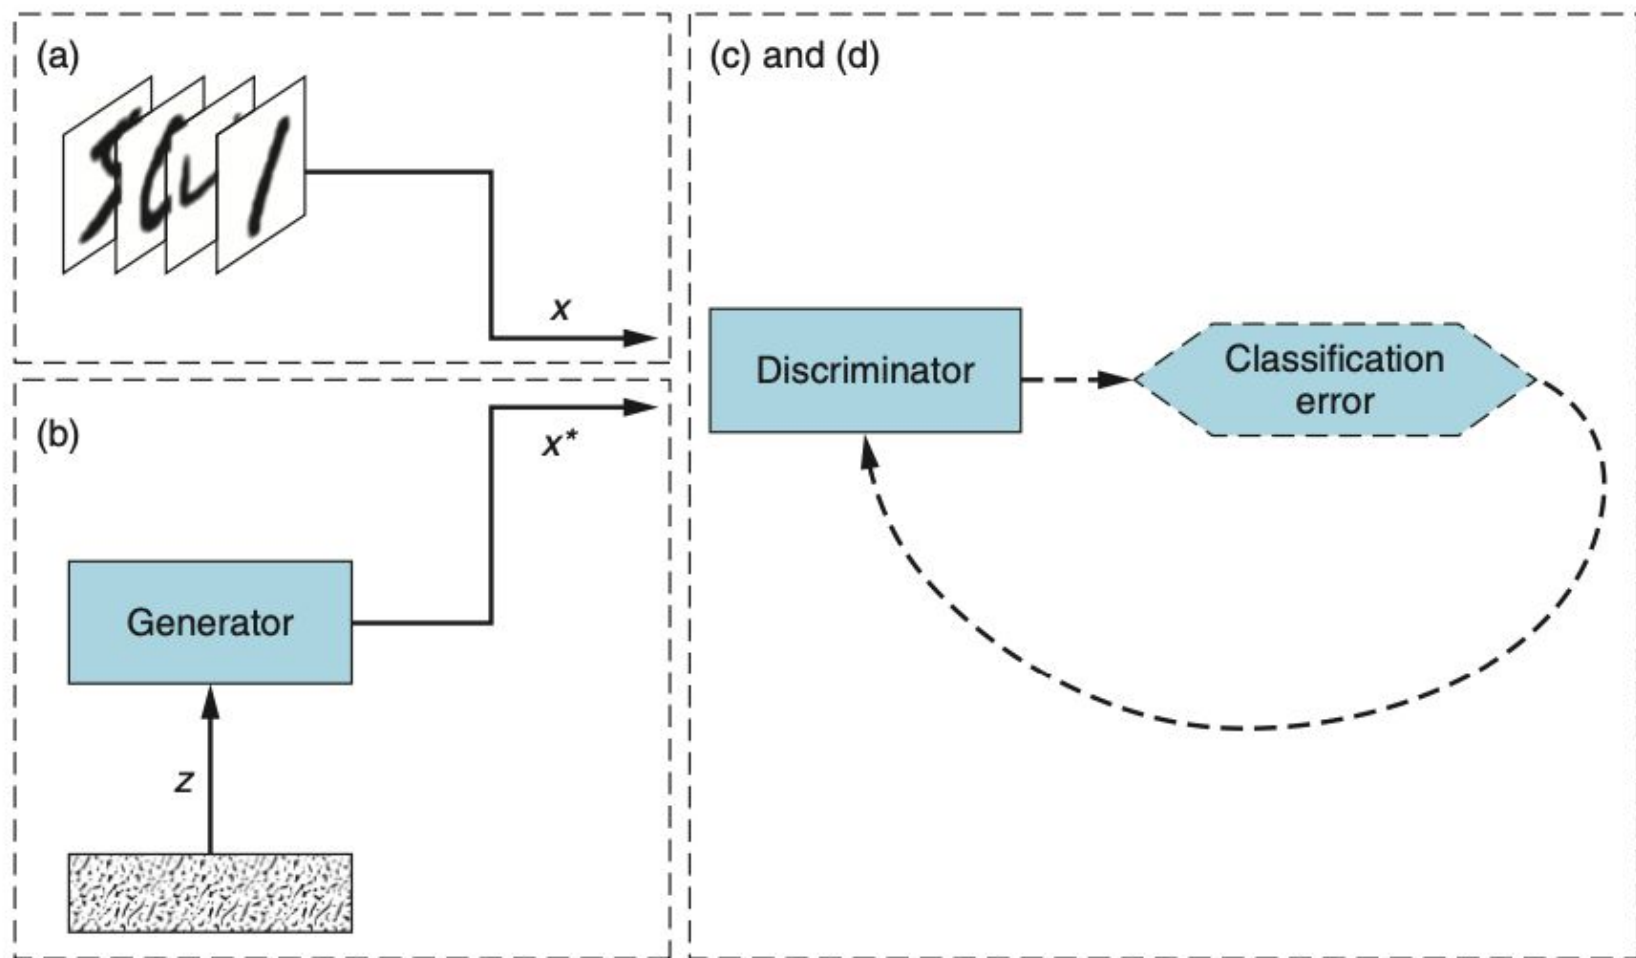


2017

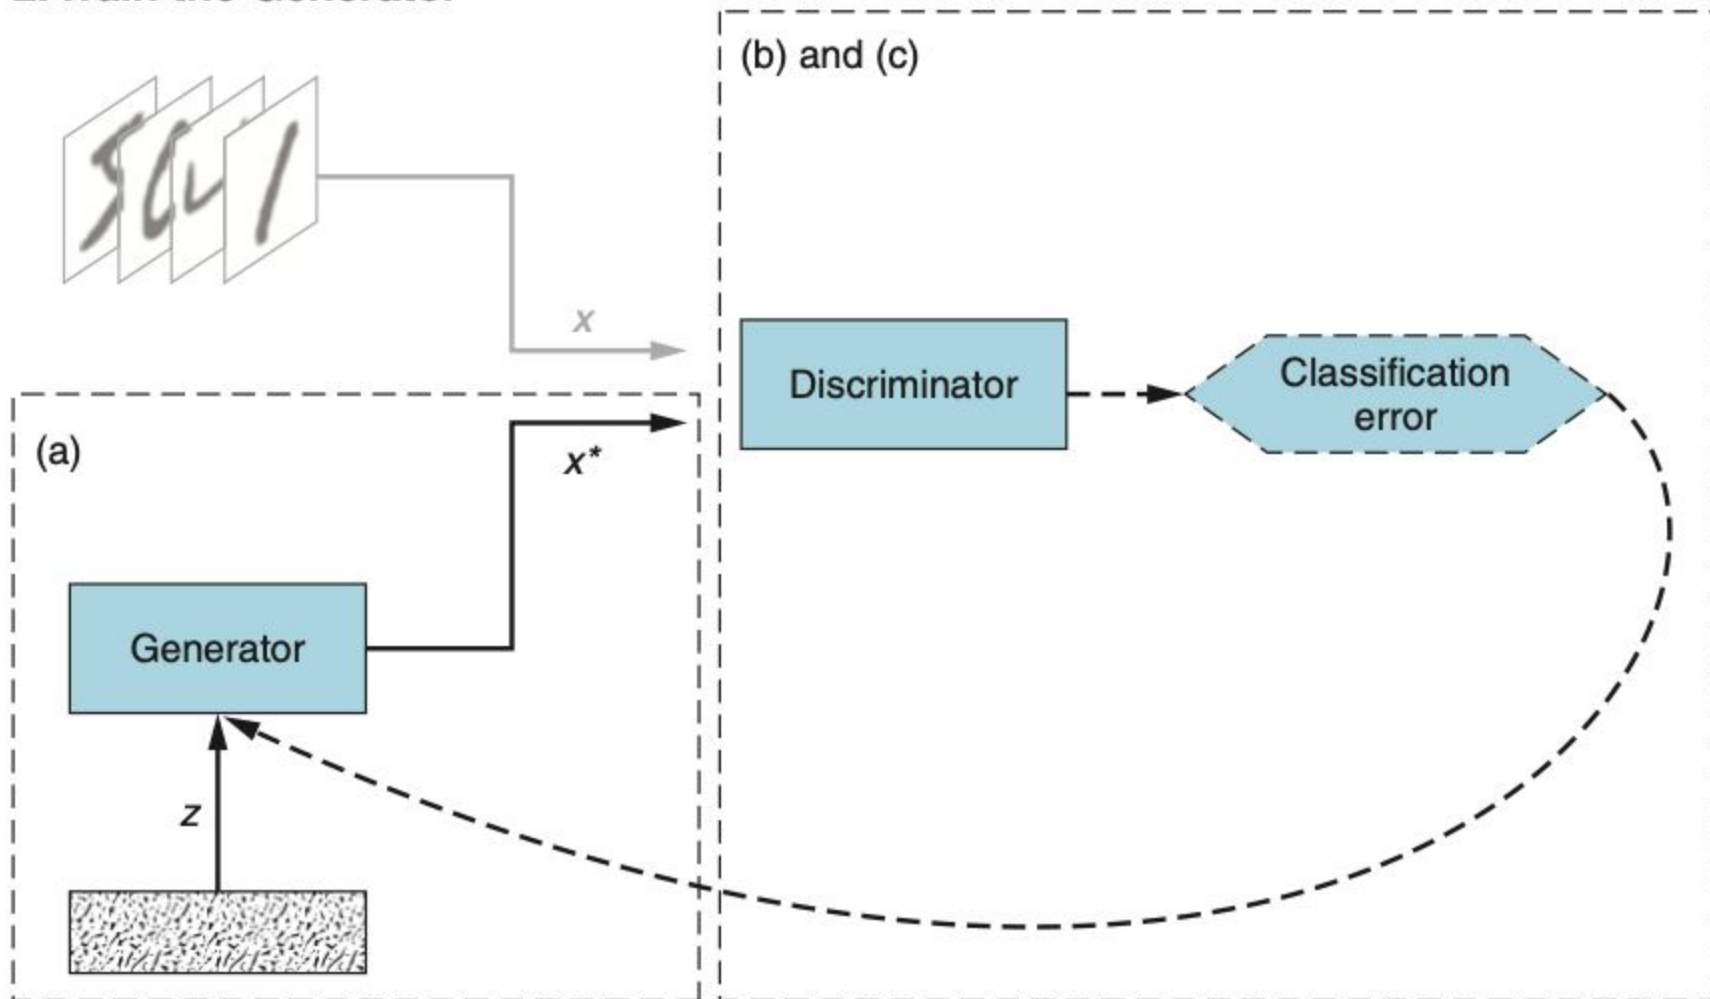
Figure 1.1 Progress in human face generation

(Source: "The Malicious Use of Artificial Intelligence: Forecasting, Prevention, and Mitigation," by Miles Brundage et al., 2018, <https://arxiv.org/abs/1802.07228>.)

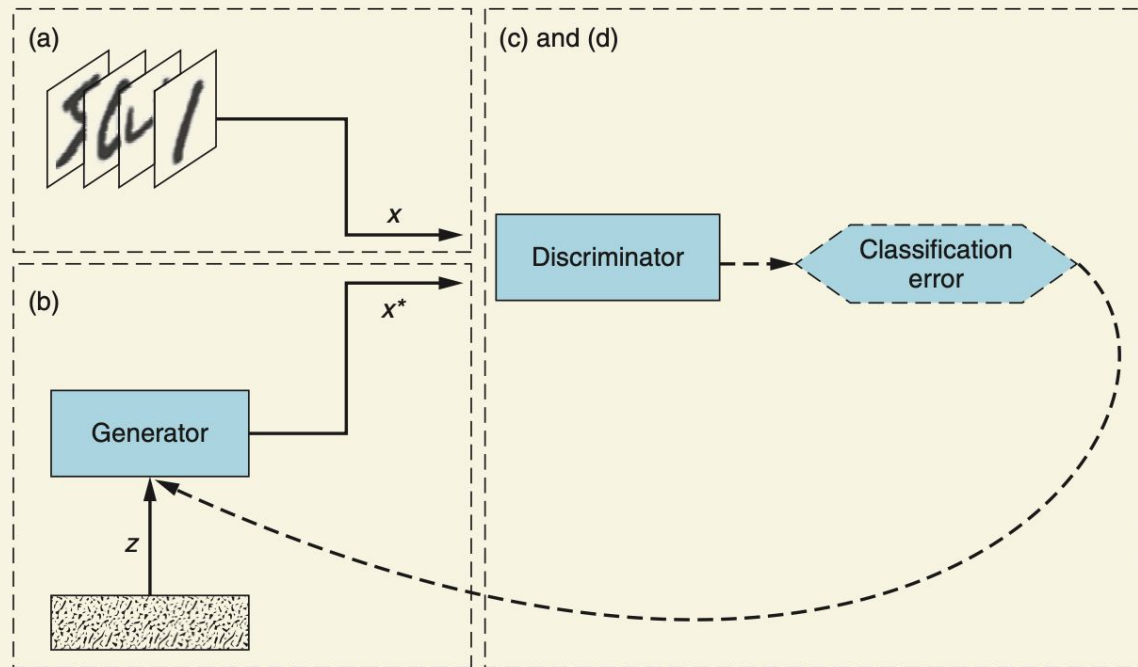
1. Train the Discriminator



2. Train the Generator



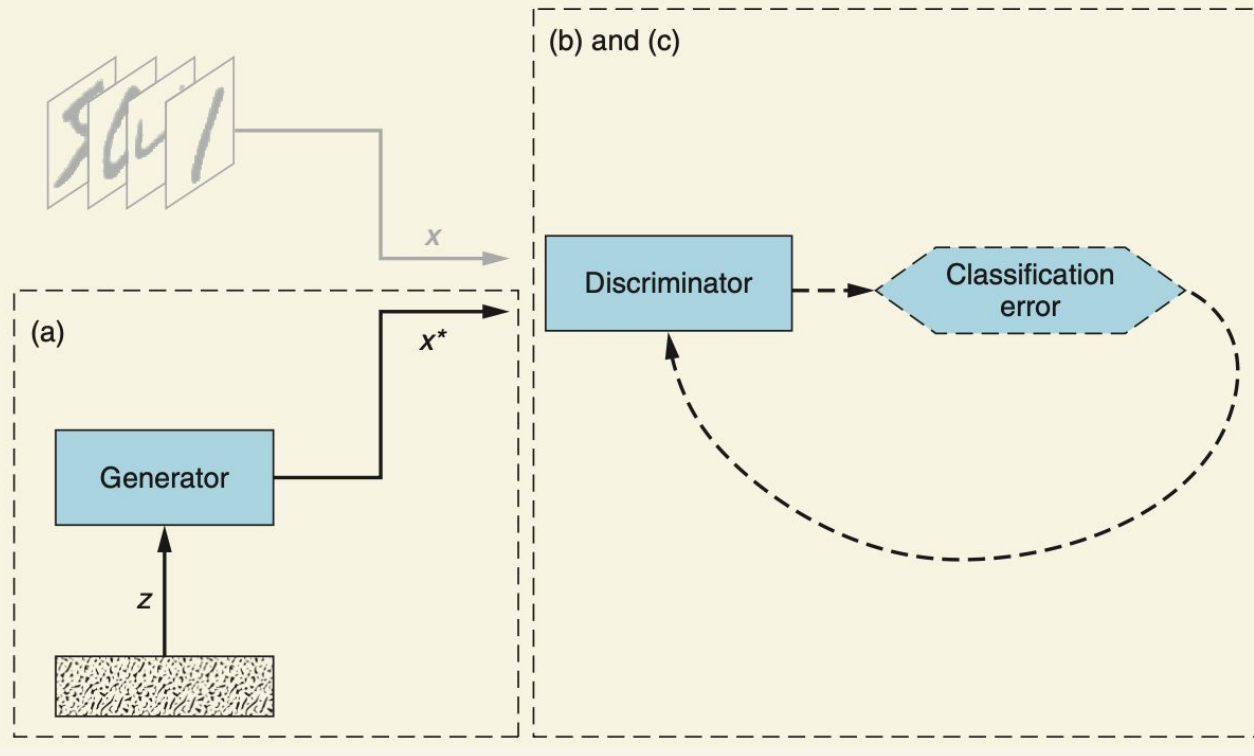
1 Train the Discriminator:



- a Take a random real example x from the training dataset.
- b Get a new random noise vector z and, using the Generator network, synthesize a fake example x^* .
- c Use the Discriminator network to classify x and x^* .
- d Compute the classification errors and backpropagate the total error to update the Discriminator weights and biases, seeking to *minimize* the classification errors.

2 Train the Generator:

2 Train the Generator:



- a** Get a new random noise vector z and, using the Generator network, synthesize a fake example x^* .
- b** Use the Discriminator network to classify x^* .
- c** Compute the classification error and backpropagate the error to update the Generator weights and biases, seeking to *maximize* the Discriminator's error.

Quiz Time!

What is the main goal of the Generator in a GAN?

- a) To detect fake data.
- b) To create fake data that looks real.
- c) To classify data into categories.
- d) To stop the Discriminator from learning.

Quiz Time!

Why do GANs improve over time during training?

- a) They use more layers in the neural network.
- b) Both Generator and Discriminator learn from each other's mistakes.
- c) The learning rate keeps increasing.
- d) The model memorizes the training set.

Quiz Time!

Which of the following is NOT a common application of GANs?

- a) Image generation
- b) Image super-resolution
- c) Predicting stock prices
- d) Deepfake video creation

Quiz Time!

True or False:

In an ideal GAN training process, the Generator becomes so good that the Discriminator can no longer tell real from fake better than random guessing.